
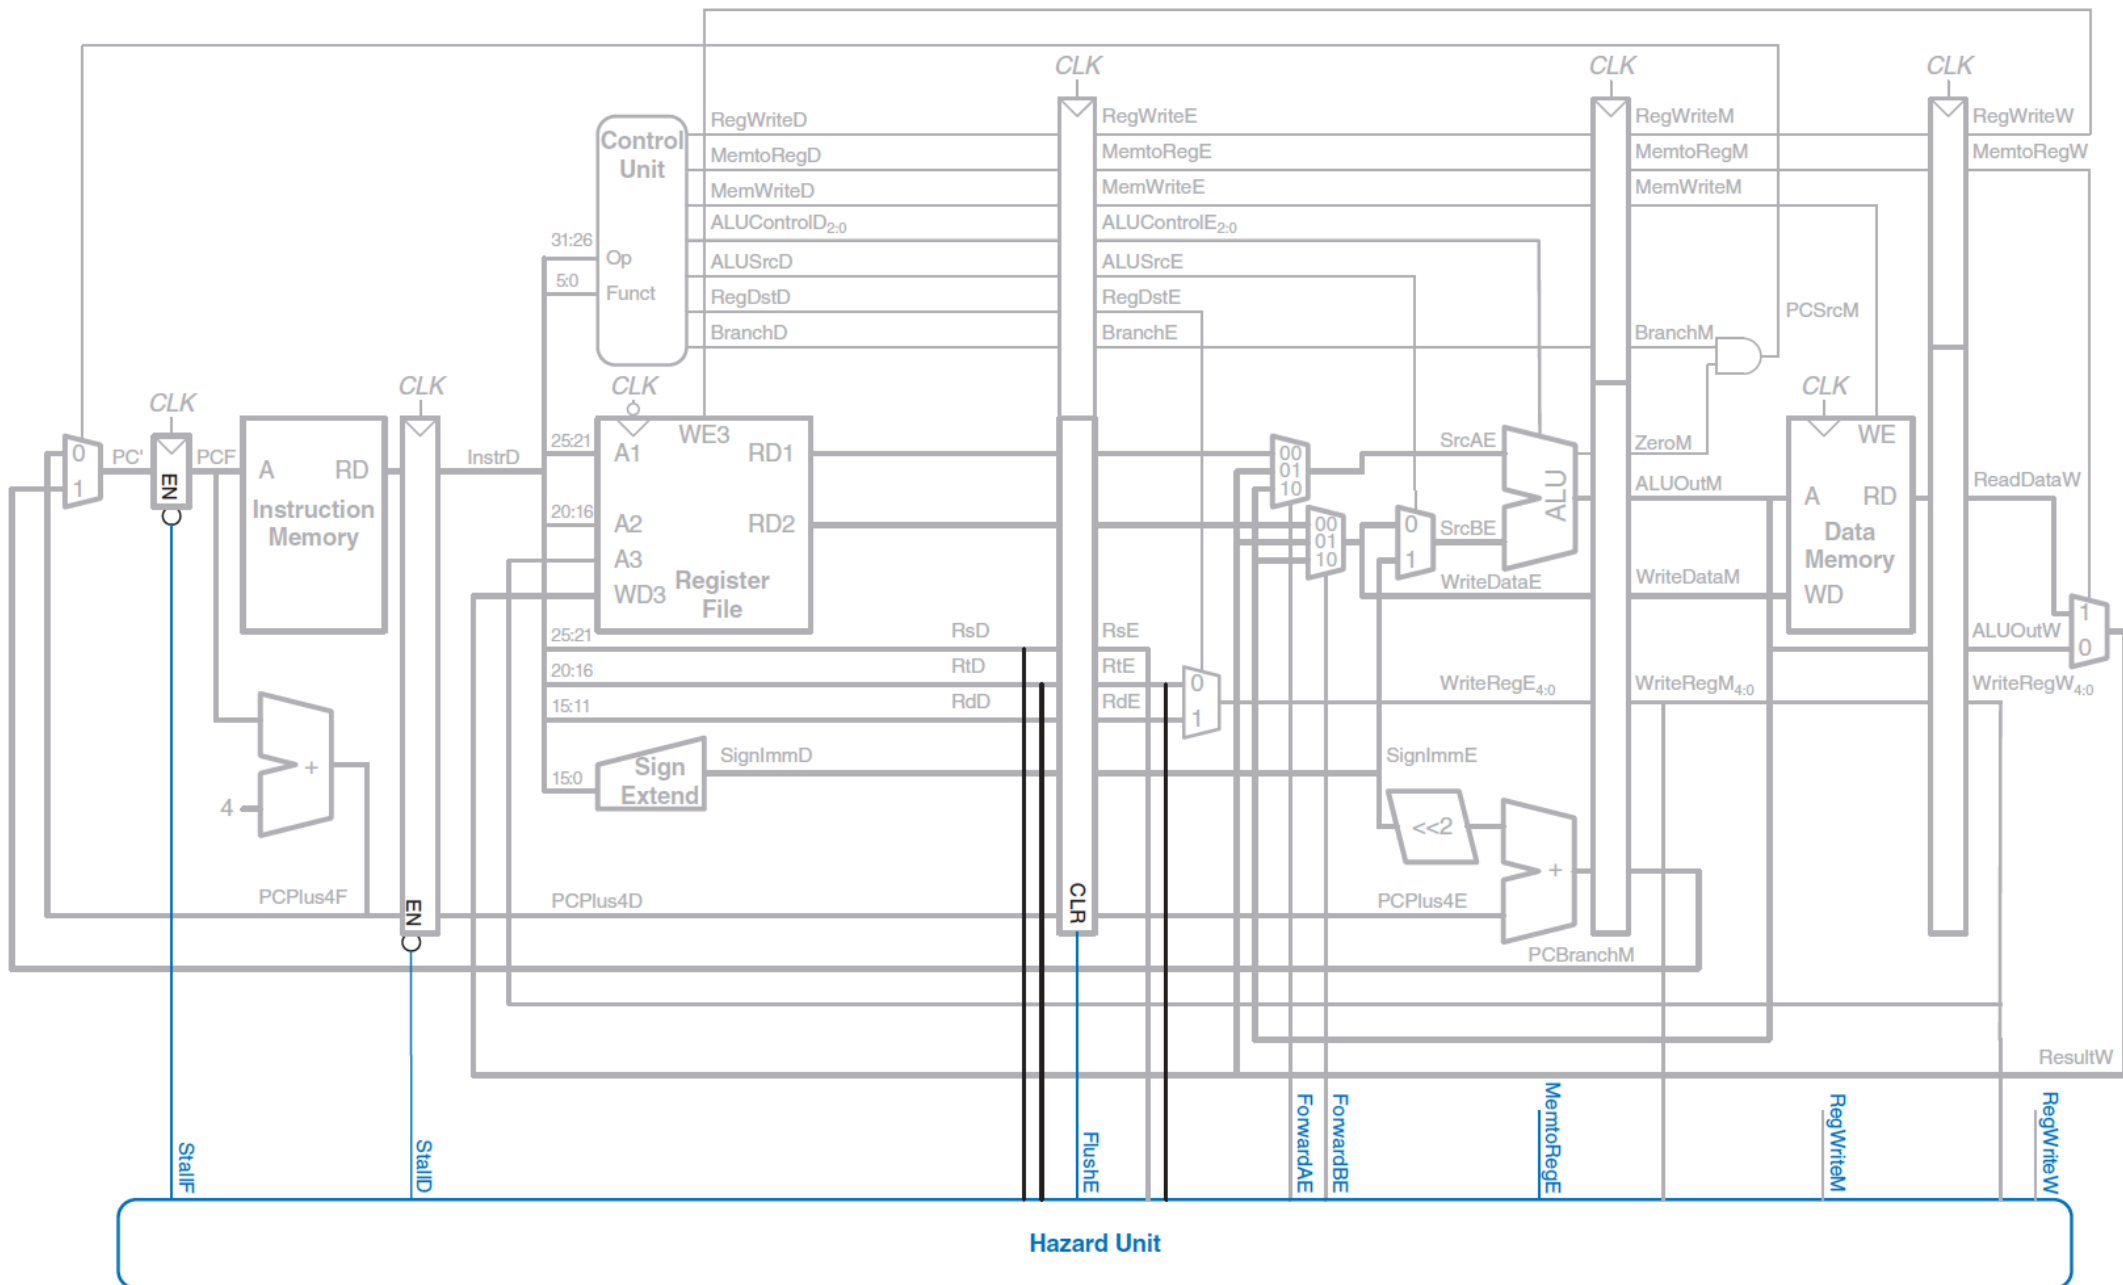


Pipeline Control Hazards





Hazards can be Data or Control

Data hazard

Instruction tries to read a register that has not been written back yet by a previous instruction.

Control hazard

Decision of what instruction to fetch next has not been made by the time the fetch takes place.

Need to enhance our pipelined processor to detect and handle hazards.

Solving Control Hazards

Processor does not know what instruction to fetch next

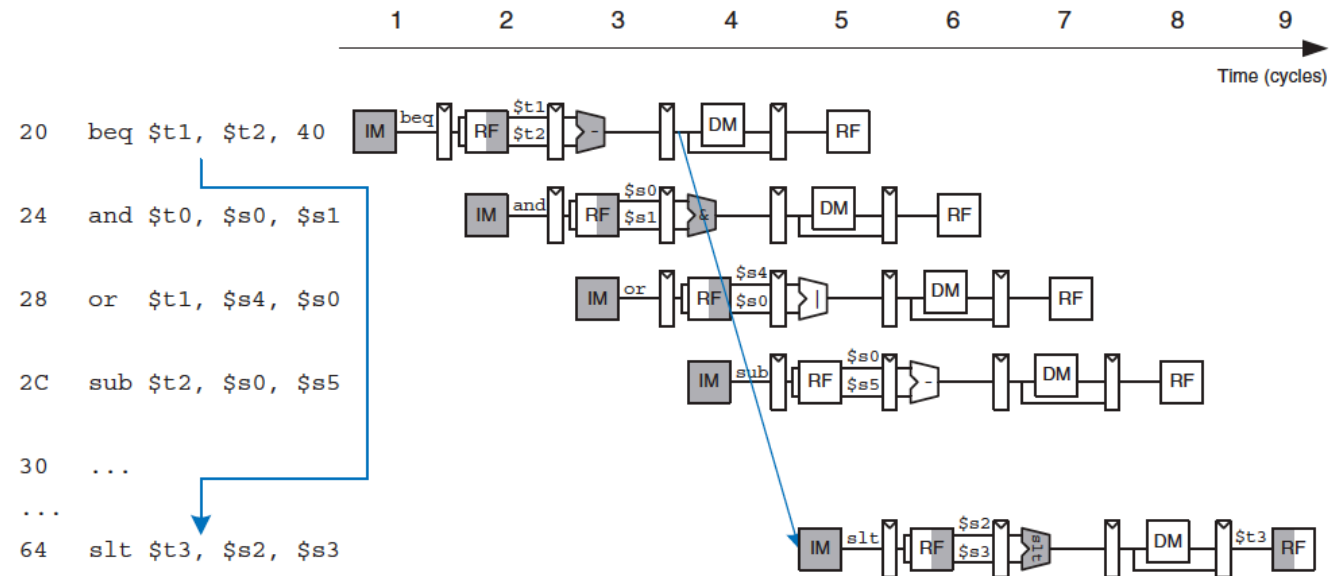
Branch decision has not been made by the time the next instruction is fetched

Can stall the pipeline until the branch decision is made

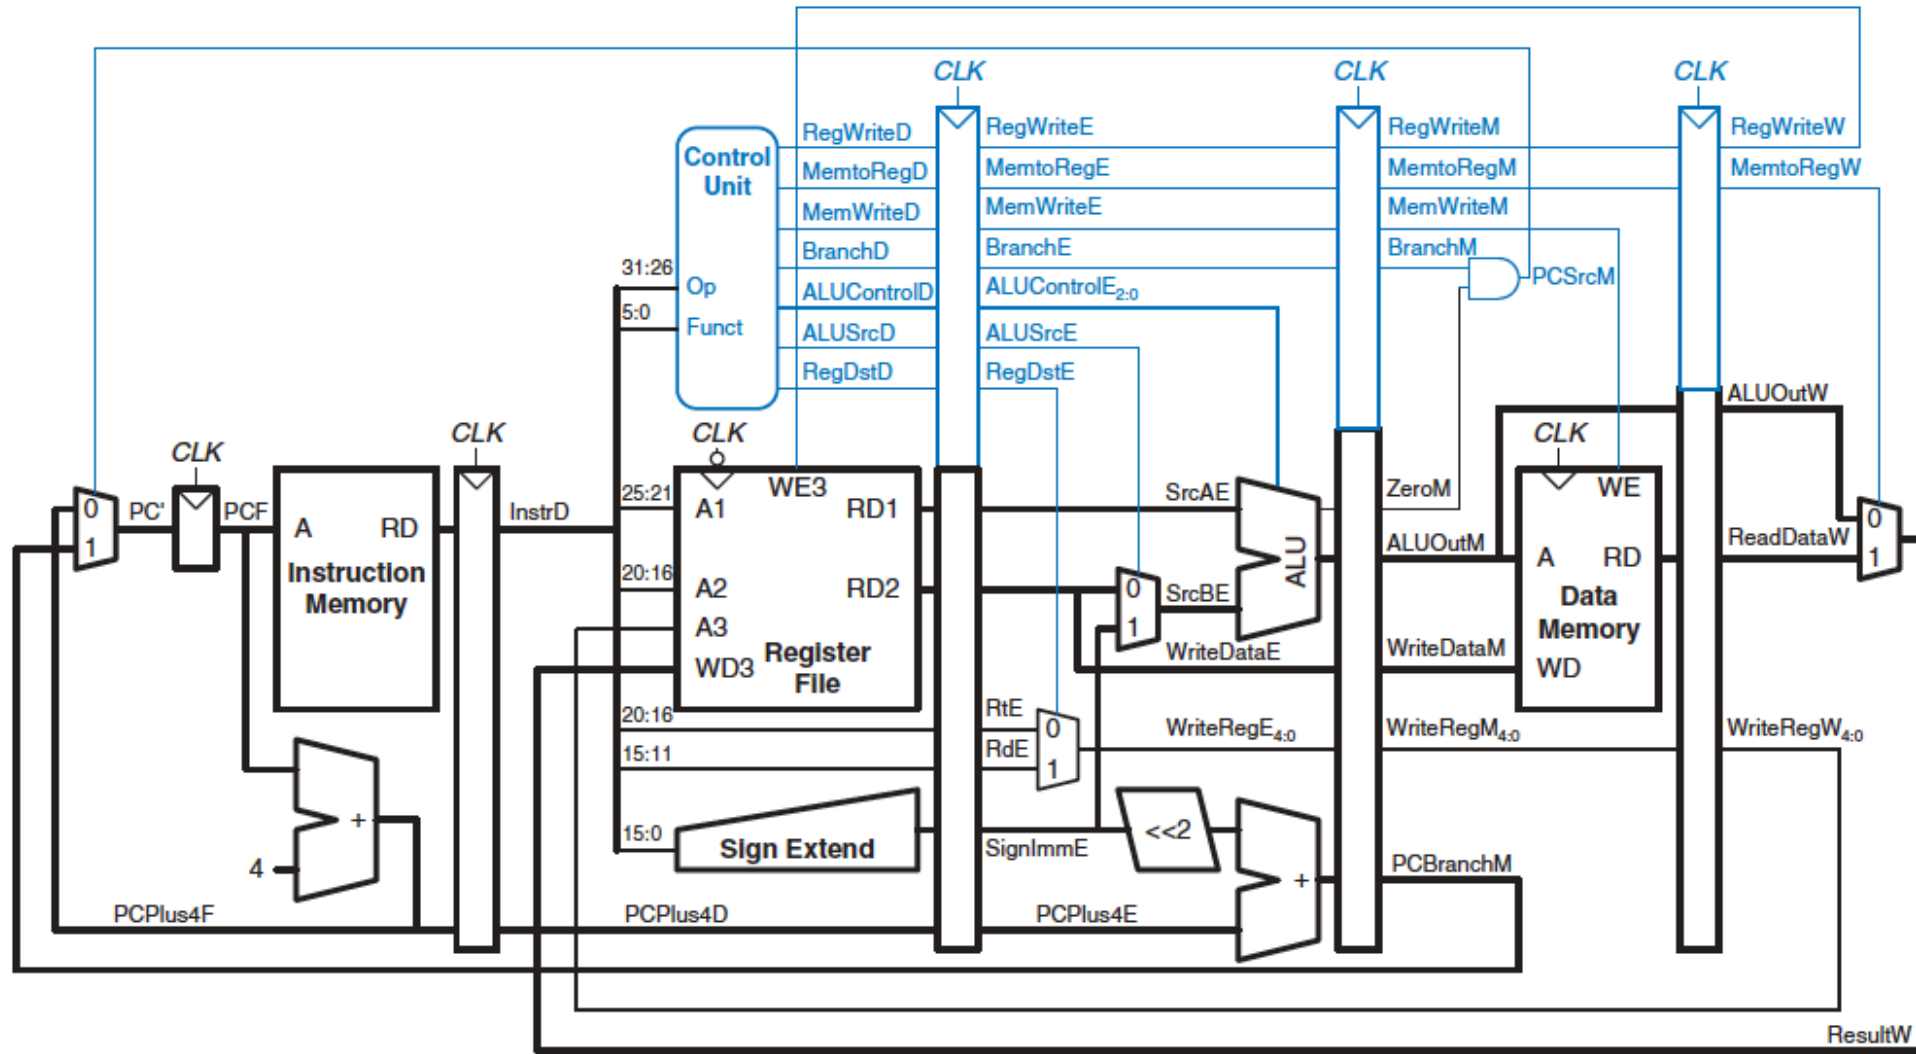
PCSrc computed

Decision is made in the Memory stage, meaning pipeline would need to stall three cycles at every branch

Drastically hamper system performance



Can you think of an alternative?



Branch Prediction

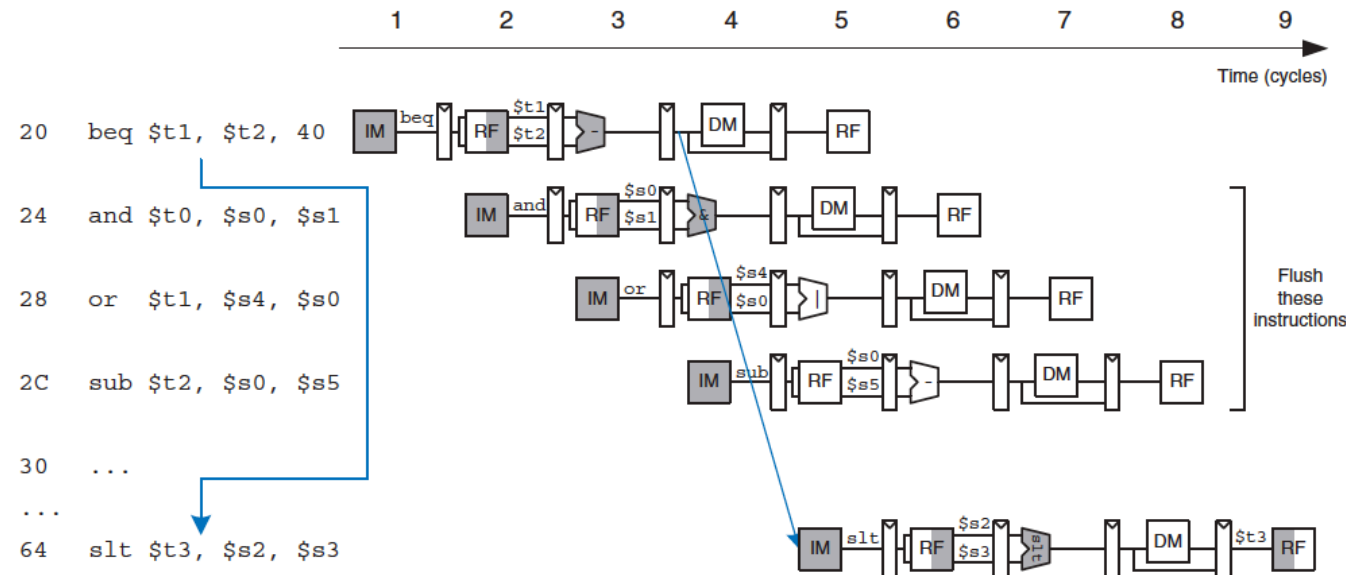
Predict whether the branch will be taken.

Act on a prediction and fill the pipeline accordingly.

If wrong, throw out the instructions that were loaded.

Flushing

Wasted instruction cycles are called the branch misprediction penalty.



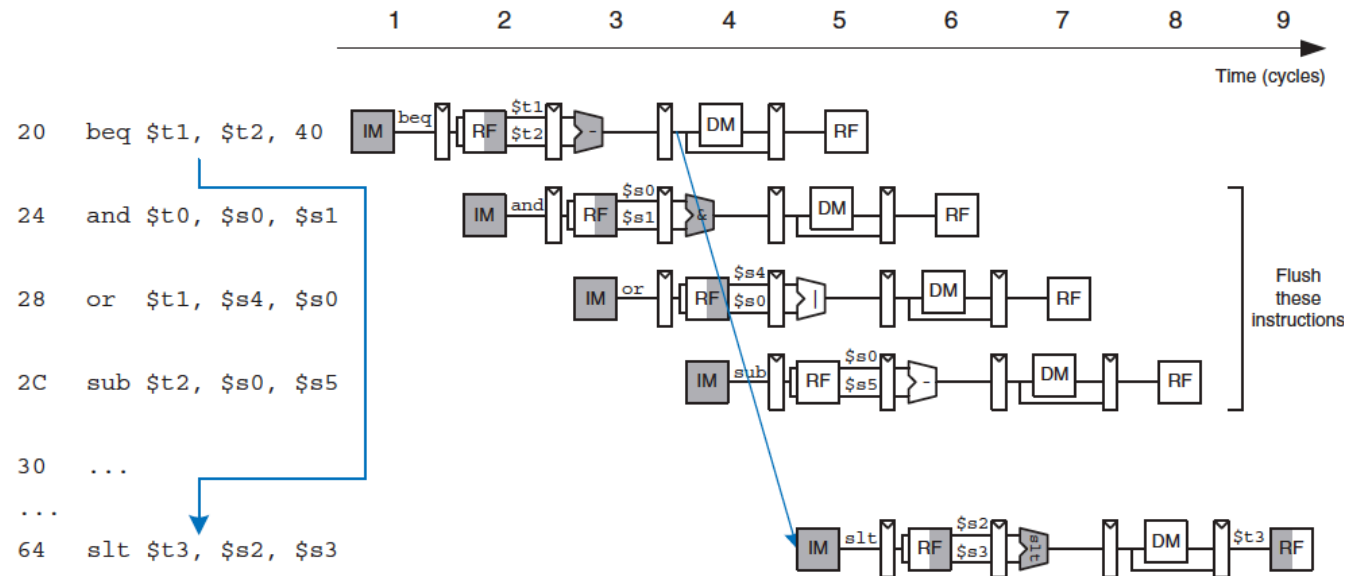
Branch Prediction

Branch from address 20 to 64 is taken

Decision not made until cycle 4, and, or and sub have already been fetched

They must be flushed

Slt fetched from line 64 on cycle 5

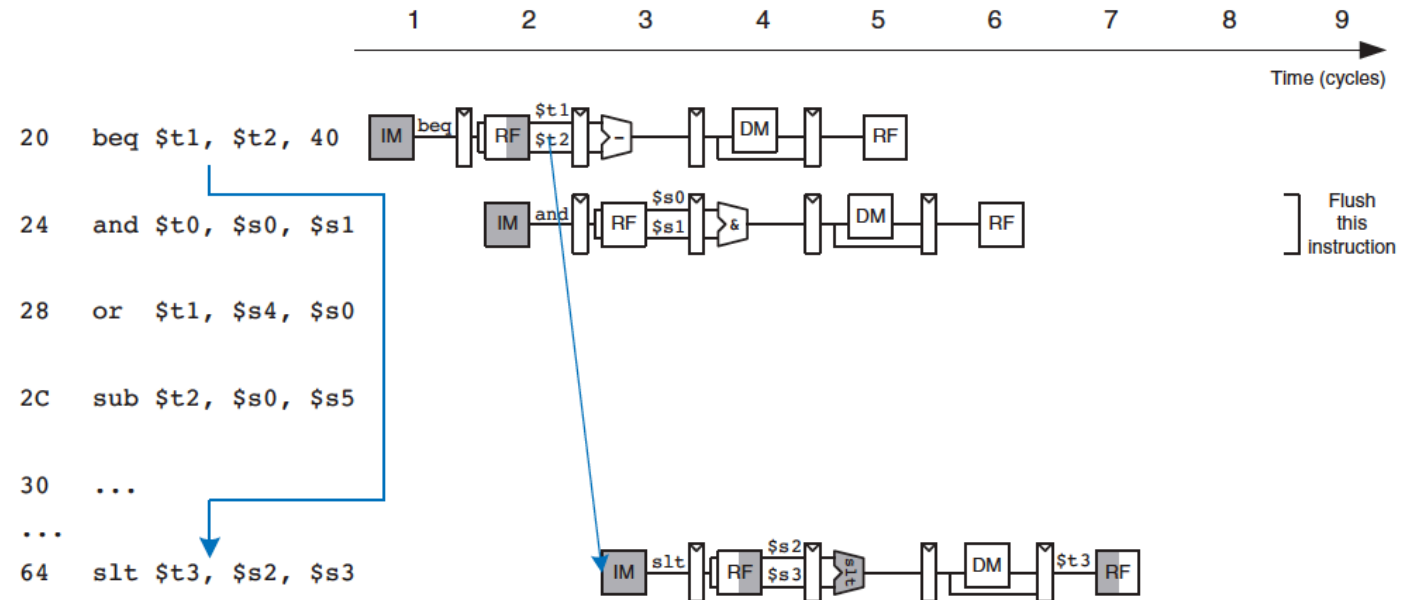


Branch Prediction

Could make our predictions earlier

Compare the value of two registers with dedicated equality comparator
Early branch detection

Branch misprediction reduced to one instruction rather than three.



CPU Modifications

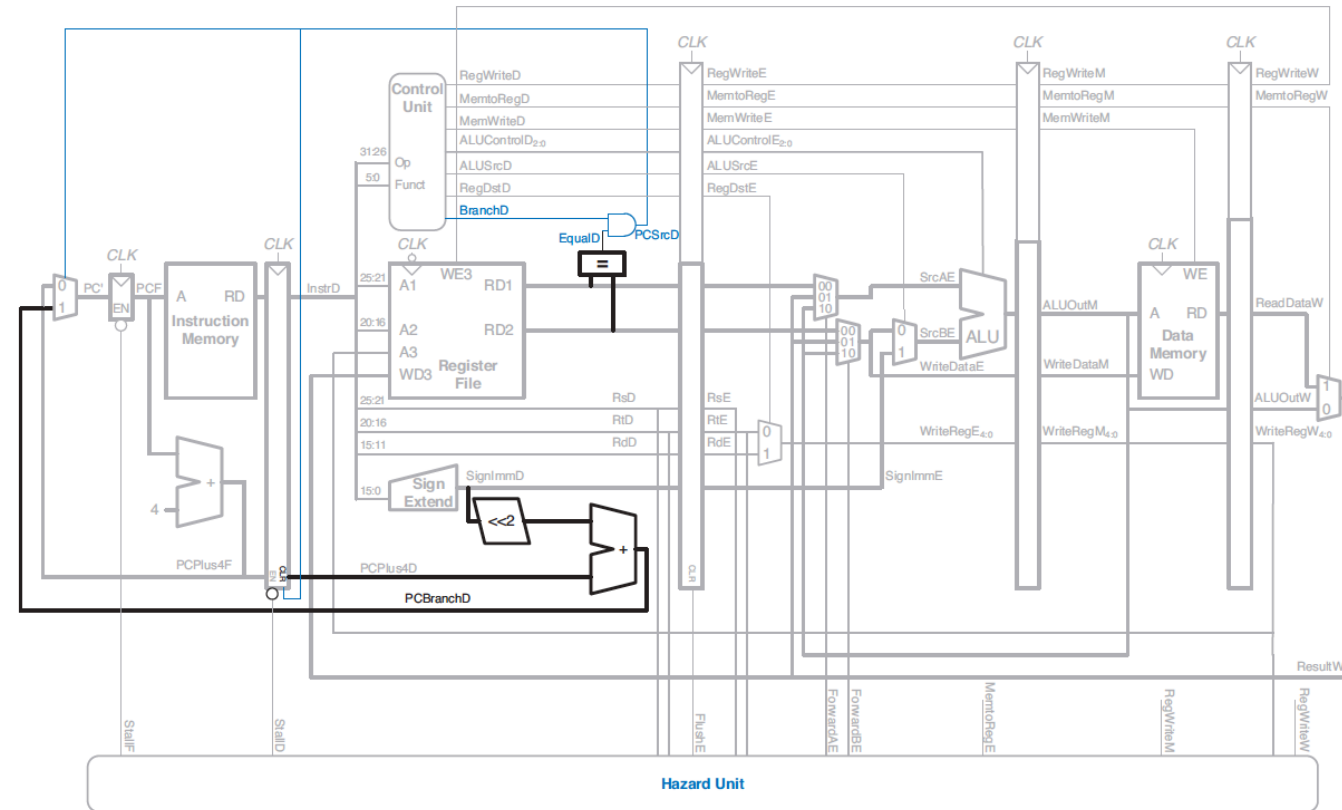
CPU modification with branch decision moved up and control hazards handled.

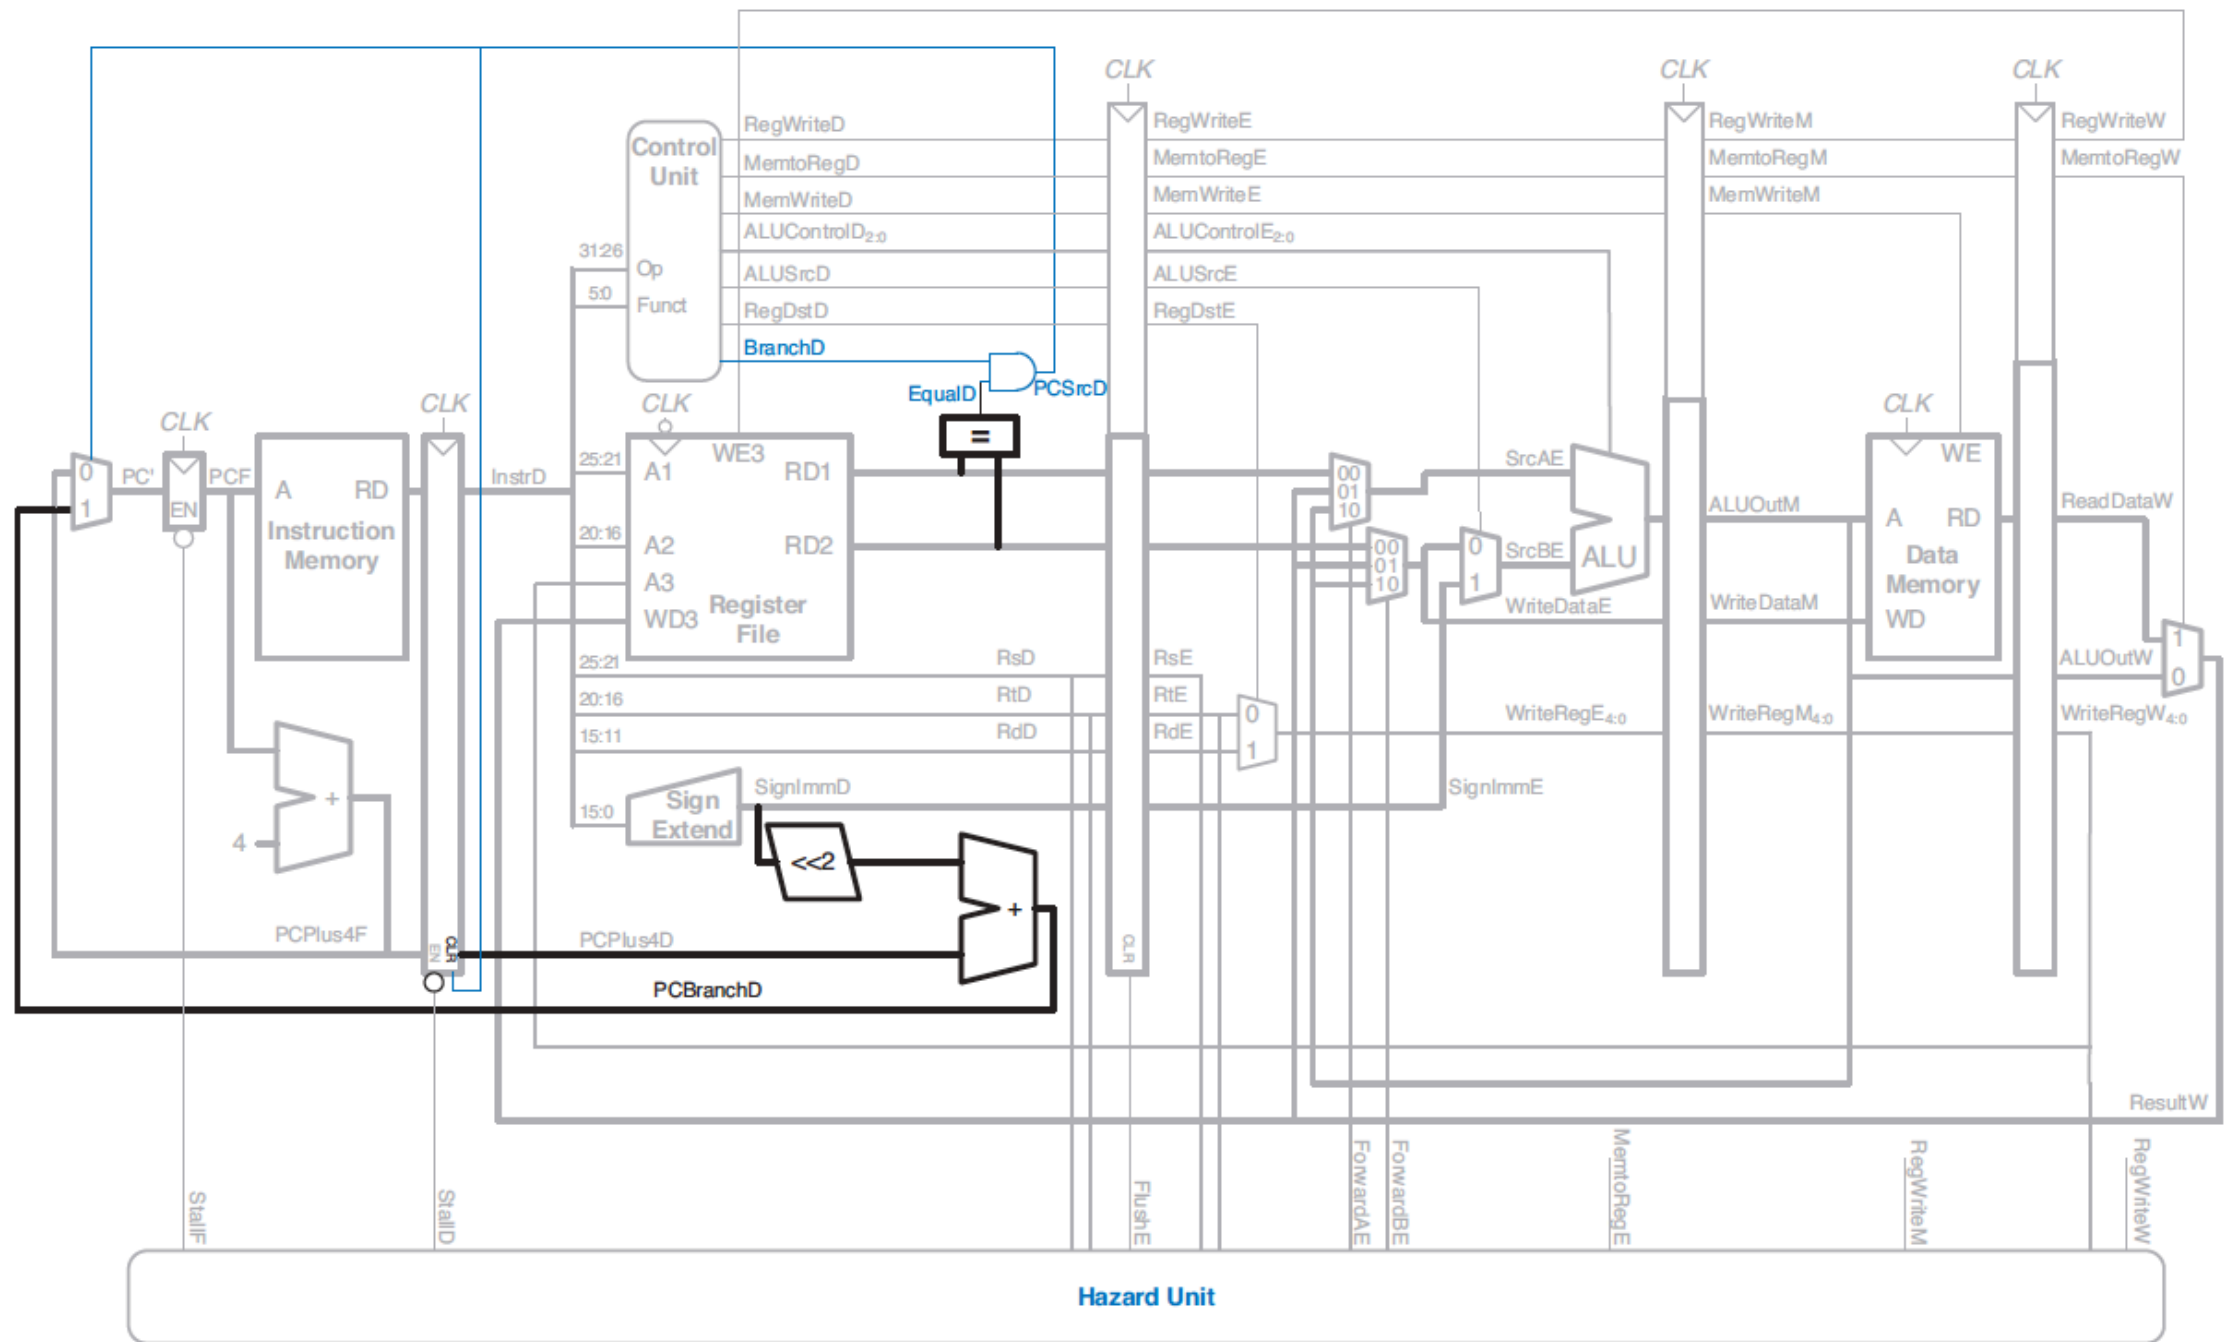
Equality comparator added to the decode stage and PCSrc AND gate is moved earlier

Determine PCSrc in the Decode rather than Memory stage

PCBranch adder must also be moved into the Decode to calculate the destination address in time.

Clear connected to PCSrcD to flush when a branch is taken.





How well can branch prediction work?

- Seems like a 50/50 chance of guessing correctly
- In reality, can do much better than this. Why? (Think about different cases where we branch)

How well can branch prediction work?

loop:

 slti \$t1, \$t0, 500

 beq \$t1, \$0, done

 ... # do work

 addi \$t0, \$t0, 1

 j loop

done:

...

Summary

1. RAW data hazards occur when an instruction depends on the result of another instruction that has not yet been written to the register file.

2. Data hazards resolved by forwarding if result is computed soon enough.

Otherwise stall pipeline until the result is available

3. Control hazards occur when the decision of what instruction to fetch has not been made by the time the next instruction must be fetched.

Solved by predicting what instruction should be fetched

Flush if prediction turns out wrong.

Moving the decision earlier reduces how many instructions must be flushed.